

# Bank Marketing MLOps Demo (Direct-Use API)

End-to-end classification workflow that predicts whether a banking customer will subscribe to a term deposit.

This fork removes the login flow: the predictor UI is available at `/app`, and the `/predict` API can be called directly.

The API **engineers helper features at inference time** (so your JSON only needs the raw fields).

---

## Tech Stack

- Python 3.11+
  - Pandas, NumPy, scikit-learn
  - MLflow (latest)
  - FastAPI + Uvicorn (served with Gunicorn in Azure)
  - Azure App Service (Linux, Basic B1)
- 

## Project Layout (high level)

```
data/           # raw dataset provided
models/         # trained MLflow model (created by training)
dist/           # packaged model zip (created by packaging)
src/            # config, data, features, training pipeline
app/            # FastAPI app (main.py + templates)
scripts/        # CLI: train, package, setup_model, run_api
azure/          # Dockerfile, entrypoint, startup.txt, deploy scripts
requirements.txt
```

Key bits: - `app/main.py` exposes: - `GET /` → marketing landing page - `GET /app` → predictor dashboard (no auth) - `POST /predict` → scoring endpoint (Pydantic schema below) - `src/features/engineering.py` defines engineered columns used by the model.

The API **applies this automatically** before prediction. - `scripts/setup_model.py` downloads/extracts a model when `MODEL_URI` is set (HTTP/HTTPS/Azure Blob/local path/zip).

---

## Local: setup, train, package, serve

```
# 1) env
python -m venv .venv
.\.venv\Scripts\Activate.ps1
```

```
pip install --upgrade pip
pip install -r requirements.txt

# 2) train (logs to ./mlruns; saves model to ./models/bank_marketing_model)
python -m scripts.train_model --experiment-name bank_marketing_classification

# 3) package (creates dist/bank_marketing_model.zip)
python -m scripts.package_model --output-dir dist

# 4) serve locally (expects model in ./models/bank_marketing_model)
uvicorn app.main:app --reload

# UI
#   http://127.0.0.1:8000/app
# API docs
#   http://127.0.0.1:8000/docs
```

JSON request for `/predict`:

```
{
  "age": 42,
  "job": "blue-collar",
  "marital": "married",
  "education": "secondary",
  "default": "no",
  "balance": 1500,
  "housing": "yes",
  "loan": "no",
  "contact": "cellular",
  "day": 15,
  "month": "may",
  "duration": 180,
  "campaign": 2,
  "pdays": 999,
  "previous": 0,
  "poutcome": "unknown"
}
```

The service engineers these features internally: `log_duration`, `log_campaign`, `is_balance_positive`, `has_previous_contact`.

# Deployment (CI/CD with GitHub → Azure App Service)

Replaces the previous “Quick Deploy”/zip-deploy section. This guide sets up **source-based CI/CD**: push to GitHub → build & deploy to your Azure Web App via **GitHub Actions** created by App Service **Deployment Center**.

## Prerequisites

- Azure subscription + **Azure CLI** installed and logged in: `az login`
- GitHub account
- A Web App (Linux) and Storage account for hosting the model (create below)
- Local repo checkout of this project

**Defaults used by scripts (change as needed):**

```
$rg = "rg-bank-marketing"  
$st = "stbankmarketing"  
$app = "bank-marketing-api"  
$container = "models"
```

---

## 0) One-time Azure resources

From the repo root, run the provisioning script to create: Resource Group, Linux Plan, Web App, Storage account + container.

```
powershell -ExecutionPolicy Bypass -File .\azure\deploy.ps1
```

If you created resources with different names, update the variables you use in the commands below.

---

## 1) Package the model and upload to Azure Blob Storage

1. **Package the trained model** (if not already):

```
python -m scripts.package_model --output-dir dist
```

This creates: `dist/bank_marketing_model.zip`.

2. **Get a connection string** and **upload** the model to the `models` container:

```
$CONN = az storage account show-connection-string `
  --name $st --resource-group $rg --query connectionString -o tsv

az storage blob upload `
  --connection-string $CONN `
  --container-name $container `
  --file dist\bank_marketing_model.zip `
  --name bank_marketing_model.zip
```

3. **Create a short-lived SAS** URL for the model:

```
$SAS = az storage blob generate-sas `
  --connection-string $CONN `
  --container-name $container `
  --name bank_marketing_model.zip `
  --permissions r --expiry (Get-Date).AddDays(7).ToString("yyyy-MM-dd") `
  -o tsv

$MODEL_URL = "https://$st.blob.core.windows.net/$container/
bank_marketing_model.zip?$SAS"
```

You can later rotate this SAS and update App Settings without changing the pipeline.

---

## 2) Configure App Settings & Startup (once)

Set the required environment variables on the Web App and define the startup command:

```
az webapp config appsettings set -g $rg -n $app --settings `
  MODEL_URI="$MODEL_URL" `
  WEBSITES_PORT=8000 `
  SCM_DO_BUILD_DURING_DEPLOYMENT=true `
  MLFLOW_TRACKING_URI="file:/home/site/wwwroot/mlruns"

# Set Startup Command so App Service uses gunicorn via our script
az webapp config set -g $rg -n $app --startup-file "bash azure/startup.txt"
```

**Why startup.txt?** It ensures dependencies are installed, the model is fetched via `scripts/setup_model.py`, and Gunicorn/Uvicorn is launched consistently on Linux App Service.

---

### 3) Push local code to a new GitHub repository

1. Initialize and commit (if you haven't already):

```
git init
git add .
git commit -m "Initial commit: bank marketing API"
```

2. Create a GitHub repo (via GitHub UI or CLI) and add it as `origin`:

```
git remote add origin https://github.com/<your-org-or-user>/<your-repo>.git
git branch -M main
git push -u origin main
```

#### If your local is behind GitHub (fast-forward safely)

Before pushing, bring your branch up to date without merging remote changes into a new merge commit:

```
git pull --rebase --autostash origin main
# autostash saves your uncommitted work, rebases your commits on top of origin/
main,
# then reapplies your local changes. Resolve conflicts if prompted.

git push origin main
```

**Notes** - Use this when you want a **linear history** (no merge commits from `git pull`). - If you already pushed your branch to GitHub and others pulled it, avoid rebasing unless everyone agrees. - If you prefer a merge-based workflow, you can instead do: `git pull origin main` (no `--rebase`).

---

### 4) Enable CI/CD from Azure App Service to GitHub (Deployment Center) from Azure App Service to GitHub (Deployment Center)

1. Open **Azure Portal** → **App Services** → *your app* → **Deployment Center**.
2. **Source**: choose **GitHub**. Authorize if prompted.
3. Select your **Organization** → **Repository** → **Branch** (e.g., `main`).
4. **Build provider**: choose **GitHub Actions**.
5. **Runtime stack**: **Python 3.11**; OS: **Linux**.
6. Click **Save** / **Finish**. Azure will:
7. Create a workflow file in your repo: `.github/workflows/<something>.yaml`.
8. Create a GitHub Secret for the publish profile and reference it in the workflow.
9. Trigger the **first CI/CD run** automatically.

If you prefer to add a workflow yourself, use the template below.

## 5) GitHub Actions workflow (auto-generated by Azure App Service Deployment Center)

When you enable CI/CD from **Azure App Service** → **Deployment Center** and select GitHub as the source with **GitHub Actions** as the build provider, Azure **auto-generates** a workflow file in your repository under `.github/workflows/`. A typical workflow looks like this (generated by Deployment Center):

```
# Docs for the Azure Web Apps Deploy action: https://github.com/Azure/webapps-deploy
# More GitHub Actions for Azure: https://github.com/Azure/actions
# More info on Python, GitHub Actions, and Azure App Service: https://aka.ms/python-webapps-actions

name: Build and deploy Python app to Azure Web App - bank-marketing-api

on:
  push:
    branches:
      - main
  workflow_dispatch:

jobs:
  build:
    runs-on: ubuntu-latest
    permissions:
      contents: read #This is required for actions/checkout

    steps:
      - uses: actions/checkout@v4

      - name: Set up Python version
        uses: actions/setup-python@v5
        with:
          python-version: '3.11'

      # 🛠️ Local Build Section (Optional)

      # The following section in your workflow is designed to catch build issues early
      # on the client side, before deployment. This can be helpful for debugging and
      # validation. However, if this step significantly increases deployment time and
      # early detection is not critical for your workflow, you may remove this section
      # to streamline the deployment process.
```

```
- name: Create and Start virtual environment and Install dependencies
  run: |
    python -m venv antenv
    source antenv/bin/activate
    pip install -r requirements.txt
```

# By default, when you enable GitHub CI/CD integration through the Azure portal, the platform automatically sets the `SCM_DO_BUILD_DURING_DEPLOYMENT` application setting to true. This triggers the use of Oryx, a build engine that handles application compilation and dependency installation (e.g., `pip install`) directly on the platform during deployment. Hence, we exclude the antenv virtual environment directory from the deployment artifact to reduce the payload size.

```
- name: Upload artifact for deployment jobs
  uses: actions/upload-artifact@v4
  with:
    name: python-app
    path: |
      .
      !antenv/
```

# 🚫 Opting Out of Oryx Build

# If you prefer to disable the Oryx build process during deployment, follow these steps:

# 1. Remove the `SCM_DO_BUILD_DURING_DEPLOYMENT` app setting from your Azure App Service Environment variables.

# 2. Refer to sample workflows for alternative deployment strategies:  
<https://github.com/Azure/actions-workflow-samples/tree/master/AppService>

```
deploy:
  runs-on: ubuntu-latest
  needs: build
  permissions:
    id-token: write #This is required for requesting the JWT
    contents: read #This is required for actions/checkout
```

```
steps:
  - name: Download artifact from build job
    uses: actions/download-artifact@v4
    with:
      name: python-app
```

```
- name: Login to Azure
  uses: azure/login@v2
  with:
```

```
    client-id: $
```

```
{{ secrets.AZUREAPPSERVICE_CLIENTID_599F19C6F97C4EF5B8D9F6CF27266AA9 }}
```

```

    tenant-id: $
  {{ secrets.AZUREAPPSERVICE_TENANTID_34AE7BFA6D9C458F88D720290F312701 }}
    subscription-id: $
  {{ secrets.AZUREAPPSERVICE_SUBSCRIPTIONID_A2D6D447C4234986A08FA51AE2BD9170 }}

- name: 'Deploy to Azure Web App'
  uses: azure/webapps-deploy@v3
  id: deploy-to-webapp
  with:
    app-name: 'bank-marketing-api'
    slot-name: 'Production'

```

**What Azure sets up for you** - The workflow file (name may vary) is committed to your repo by the Deployment Center wizard. - **Authentication:** uses **OpenID Connect (OIDC)** via `azure/login@v2` with secrets added as GitHub Action secrets by the wizard. You don't need to create a publish profile secret for this flow. - **Build:** Oryx runs on the App Service side by default because `SCM_DO_BUILD_DURING_DEPLOYMENT=true` is set in App Settings. The local build step in the workflow is optional and can be removed for faster runs. - **Artifact:** The workflow uploads the repo as an artifact in the `build` job and downloads it in the `deploy` job.

If you later customize this workflow (e.g., add tests, linting, packaging, or multi-env deployments), keep the `azure/login` and `azure/webapps-deploy` steps intact, and ensure `app-name` matches your Web App.

## 6) Verify a deployment

After the workflow completes:

```

az webapp browse -g $rg -n $app
# UI:    https://{app}.azurewebsites.net/app
# Docs:  https://{app}.azurewebsites.net/docs

```

If you change the model location/SAS later, update **MODEL\_URI** in App Settings and **Restart** the app:

```

az webapp restart -g $rg -n $app

```

## Troubleshooting CI/CD

- **Workflow fails on deploy**
- Ensure the **publish profile secret** is present and referenced by the exact name in the YAML.
- Confirm the `app-name` in the workflow matches your Azure Web App name.



- **502/timeout after deploy**
- Confirm `WEBSITES_PORT=8000` is set and **Startup Command** = `bash azure/startup.txt`.
- First start may take time to install Python deps and download the model.
- **500 "Model not available"**
- SAS expired or `MODEL_URI` incorrect. Rotate SAS and update App Settings, then restart.
- **403 when fetching model**
- SAS read permission or container mismatch. Regenerate SAS with `--permissions r` and correct blob name.
- **500 "columns are missing" / wrong predictions**
- Ensure this fork (which **engineers features inside** `/predict`) is what you deployed.

## Configuration

Setting	Purpose	Default / Example
<code>MODEL_URI</code>	Remote/Local URI to <b>zip or directory</b> with MLflow model. Supports HTTPS & Azure Blob.	<code>https://stbankmarketing.blob.core.windows.net/models/bank_marketing_model.zip?...</code>
<code>MODEL_LOCAL_PATH</code>	Where the model is placed inside the app container.	<code>model/bank_marketing_model</code>
<code>MLFLOW_TRACKING_URI</code>	Where MLflow logs (for the app).	<code>file:/home/site/wwwroot/mlruns</code>
<code>WEBSITES_PORT</code>	Port to listen on in Azure App Service (Linux).	<code>8000</code>
<code>PYTHONPATH</code>	Optional module path.	<code>/home/site/wwwroot</code>

## API Summary

- **POST** `/predict` → returns `{{ "probability": float, "prediction": 0|1 }}`  
Request fields: `age, job, marital, education, default, balance, housing, loan, contact, day, month, duration, campaign, pdays, previous, poutcome`.
- **GET** `/app` → browser UI to submit a form and see the prediction.
- **GET** `/docs` → interactive OpenAPI docs.

---

## Notes & Next Steps

- To make `/` go straight to the predictor, change the index route to return `render_dashboard_html()` or add a redirect in `app/main.py`.
- **CI/CD:** You now deploy by **pushing to GitHub**. The workflow will build and publish to Azure automatically.  
Consider a second workflow for the offline ML loop: run `train_model` → evaluate → `package_model` → upload to Blob → update `MODEL_URI` via CLI or Portal.
- Consider storing secrets (SAS, connection strings) in Azure Key Vault and referencing them from App Settings.

---

**Dataset:** UCI Bank Marketing – for research/education.

**Security:** Never commit secrets. Rotate SAS regularly.